

# Exploratory Data Analysis - Laptops Pricing dataset

## Objectives

- Visualize individual feature patterns
- Run descriptive statistical analysis on the dataset
- Use groups and pivot tables to find the effect of categorical variables on price
- Use Pearson Correlation to measure the interdependence between variables

## Setup

For this lab, we will be using the following libraries:

- `skillsnetwork` for downloading the data
- `pandas` ([https://pandas.pydata.org/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=10006555&utm\\_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01](https://pandas.pydata.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01)) for managing the data.
- `numpy` ([https://numpy.org/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=10006555&utm\\_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01](https://numpy.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01)) for mathematical operations.
- `scipy` ([https://docs.scipy.org/doc/scipy/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=10006555&utm\\_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01](https://docs.scipy.org/doc/scipy/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01)) for statistical operations.
- `seaborn` ([https://seaborn.pydata.org/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=10006555&utm\\_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01](https://seaborn.pydata.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01)) for visualizing the data.
- `matplotlib` ([https://matplotlib.org/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=10006555&utm\\_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01](https://matplotlib.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMML0187ENSkillsNetwork31430127-2021-01-01)) for additional plotting tools.

## Install Required Libraries

### Importing Required Libraries

We import all required libraries in one place:

```
In [35]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
%matplotlib inline
```

## Import the dataset

We will use the modified pre-processed version of the data set from the last lab : `laptop_pricing_cleaned.csv`.

```
In [36]: file_name = "laptop_pricing_cleaned.csv"
```

Import the file to a pandas dataframe.

```
In [37]: df = pd.read_csv(file_name, header=0)
```

Print the first 5 entries of the dataset to confirm loading.

In [38]:

df.head(10)

Out[38]:

	Manufacturer	Category	GPU	OS	CPU_core	Screen_Size_inch	CPU_frequency	RAM_GB	Storage_GB_SSD	Weight_pounds	Price	Price-binned
0	Acer	4	2	1	5	14.0	0.551724	8	256	3.52800	978	L
1	Dell	3	1	1	3	15.6	0.689655	4	256	4.85100	634	L
2	Dell	3	1	1	7	15.6	0.931034	8	256	4.85100	946	L
3	Dell	4	2	1	5	13.3	0.551724	8	128	2.69010	1244	L
4	HP	4	2	1	7	15.6	0.620690	8	256	4.21155	837	L
5	Dell	3	1	1	5	15.6	0.551724	8	256	4.85100	1016	L
6	HP	3	3	1	5	15.6	0.551724	8	256	4.63050	1117	L
7	Acer	3	2	1	5	15.0	0.551724	4	256	4.85100	866	L
8	Dell	3	1	1	5	15.6	0.862069	4	256	5.07150	812	L
9	Acer	3	3	1	7	15.0	0.620690	8	256	4.85100	1068	L

In [39]:

df.dtypes

Out[39]:

Manufacturer                  object  
Category                     int64  
GPU                           int64  
OS                            int64  
CPU\_core                      int64  
Screen\_Size\_inch              float64  
CPU\_frequency                 float64  
RAM\_GB                        int64  
Storage\_GB\_SSD                int64  
Weight\_pounds                 float64  
Price                         int64  
Price-binned                  object  
Screen-Full\_HD                int64  
Screen-IPS\_panel              int64  
dtype: object

## Task 1 - Visualize individual feature patterns

### Continuous valued features


Generate regression plots for each of the parameters "CPU\_frequency", "Screen\_Size\_inch" and "Weight\_pounds" against "Price". Also, print the value of correlation of each feature with "Price".

In [40]:

# CPU\_frequency plot  
sns.regplot(x="CPU\_frequency", y="Price", data = df)  
plt.ylim(0,)


Out[40]:

(0.0, 3974.15)




```
In [41]: # Screen_Size_inch plot
sns.regplot( x= "Screen_Size_inch", y="Price", data=df)
plt.ylim(0,)
```

Out[41]: (0.0, 3974.15)



```
In [42]: # Weight_pounds plot
sns.regplot(x="Weight_pounds", y="Price", data=df)
plt.ylim(0,)
```

Out[42]: (0.0, 3974.15)



```
In [43]: 1 # Correlation values of the three attributes with Price
2 df_corr = df[['CPU_frequency', 'Screen_Size_inch', 'Weight_pounds', 'Price']]
3 df_corr.corr()
```

Out[43]:

	CPU_frequency	Screen_Size_inch	Weight_pounds	Price
CPU_frequency	1.000000	-0.000948	0.066522	0.366666
Screen_Size_inch	-0.000948	1.000000	0.797534	-0.110644
Weight_pounds	0.066522	0.797534	1.000000	-0.050312
Price	0.366666	-0.110644	-0.050312	1.000000

```
In [44]: # List of attributes of interest
attributes = ['CPU_frequency', 'Screen_Size_inch', 'Weight_pounds']

# Calculate correlations with 'Price'
correlations = df[attributes + ['Price']].corr()['Price'][attributes]
print(correlations)

CPU_frequency      0.366666
Screen_Size_inch   -0.110644
Weight_pounds      -0.050312
Name: Price, dtype: float64
```


Interpretation: "CPU\_frequency" has a 36% positive correlation with the price of the laptops. The other two parameters have weak correlation with price.

## Categorical features

Generate Box plots for the different feature that hold categorical values. These features would be "Category", "GPU", "OS", "CPU\_core", "RAM\_GB", "Storage\_GB\_SSD"


```
In [46]: # Category Box plot
sns.boxplot(x="Category", y="Price", data=df)
```

```
Out[46]: <Axes: xlabel='Category', ylabel='Price'>
```




```
In [47]: # GPU Box plot  
sns.boxplot(x="GPU", y="Price", data=df)
```

```
Out[47]: <Axes: xlabel='GPU', ylabel='Price'>
```




```
In [48]: # OS Box plot  
sns.boxplot(x="OS", y="Price", data=df)
```

```
Out[48]: <Axes: xlabel='OS', ylabel='Price'>
```




```
In [49]: # CPU_core Box plot
sns.boxplot(x="CPU_core", y="Price", data=df)
```

Out[49]: <Axes: xlabel='CPU\_core', ylabel='Price'>



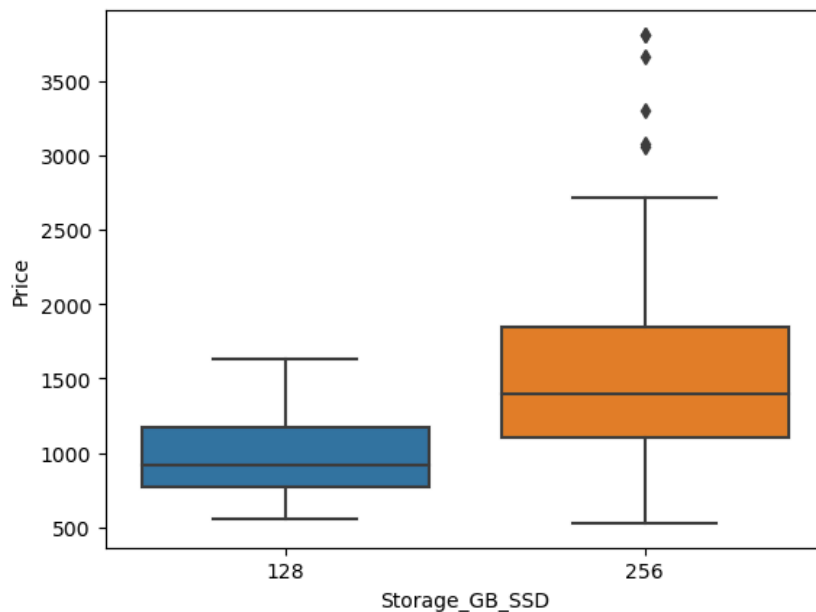
```
In [50]: # RAM_GB Box plot
sns.boxplot(x="RAM_GB", y="Price", data=df)
```

Out[50]: <Axes: xlabel='RAM\_GB', ylabel='Price'>



```
In [51]: # Storage_GB_SSD Box plot
sns.boxplot(x="Storage_GB_SSD", y="Price", data=df)
```

```
Out[51]: <Axes: xlabel='Storage_GB_SSD', ylabel='Price'>
```



## Task 2 - Descriptive Statistical Analysis

Generate the statistical description of all the features being used in the data set. Include "object" data types as well.

```
In [53]: df.describe(include=['object'])
```

```
Out[53]:
```

	Manufacturer	Price-binned
count	238	238
unique	11	3
top	Dell	Low
freq	71	160

```
In [54]: df.describe()
```

```
Out[54]:
```

	Category	GPU	OS	CPU_core	Screen_Size_inch	CPU_frequency	RAM_GB	Storage_GB_SSD	Weight_pounds	
count	238.000000	238.000000	238.000000	238.000000	238.000000	238.000000	238.000000	238.000000	238.000000	2
mean	3.205882	2.151261	1.058824	5.630252	14.688655	0.813822	7.882353	245.781513	4.106221	14
std	0.776533	0.638282	0.235790	1.241787	1.166045	0.141860	2.482603	34.765316	1.078442	5
min	1.000000	1.000000	1.000000	3.000000	12.000000	0.413793	4.000000	128.000000	1.786050	5
25%	3.000000	2.000000	1.000000	5.000000	14.000000	0.689655	8.000000	256.000000	3.246863	10
50%	3.000000	2.000000	1.000000	5.000000	15.000000	0.862069	8.000000	256.000000	4.106221	13
75%	4.000000	3.000000	1.000000	7.000000	15.600000	0.931034	8.000000	256.000000	4.851000	17
max	5.000000	3.000000	2.000000	7.000000	17.300000	1.000000	16.000000	256.000000	7.938000	38

## Task 3 - GroupBy and Pivot Tables

Group the parameters "GPU", "CPU\_core" and "Price" to make a pivot table and visualize this connection using the pcolor plot.

```
In [56]: # Create the group
df_groupe = df[["GPU", "CPU_core", "Price"]]
df_grouped = df_groupe.groupby(["GPU", "CPU_core"], as_index=False).mean()
```

```
In [58]: # Create the Pivot table
grouped_pivot = df_grouped.pivot(index="GPU", columns="CPU_core")
grouped_pivot
```

Out[58]:

Price			
CPU_core	3	5	7
GPU			
1	769.250000	998.500000	1167.941176
2	785.076923	1462.197674	1744.621622
3	784.000000	1220.680000	1945.097561

```
In [59]: # Create the Plot
fig, ax = plt.subplots()
im = ax.pcolor(grouped_pivot, cmap='RdBu')


#Label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index

#move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

fig.colorbar(im)
```

Out[59]: <matplotlib.colorbar.Colorbar at 0x2d21925e510>



## Task 4 - Pearson Correlation and p-values

Use the `scipy.stats.pearsonr()` function to evaluate the Pearson Coefficient and the p-values for each parameter tested above. This will help you determine the parameters most likely to have a strong effect on the price of the laptops.



```
In [61]: for param in ['RAM_GB', 'CPU_frequency', 'Storage_GB_SSD', 'Screen_Size_inch', 'Weight_pounds', 'CPU_core', 'OS', 'GPU',
    pearson_coef, p_value = stats.pearsonr(df[param], df['Price'])
    print(param)
    print("The Pearson Correlation Coefficient for ", param, " is", pearson_coef, " with a P-value of P =", p_value
```

RAM\_GB

The Pearson Correlation Coefficient for RAM\_GB is 0.5492972971857844 with a P-value of P = 3.681560628842868e-20

CPU\_frequency

The Pearson Correlation Coefficient for CPU\_frequency is 0.36666555832636644 with a P-value of P = 5.502463689008642e-09

Storage\_GB\_SSD

The Pearson Correlation Coefficient for Storage\_GB\_SSD is 0.2434207552181029 with a P-value of P = 0.00014898923191724174

Screen\_Size\_inch

The Pearson Correlation Coefficient for Screen\_Size\_inch is -0.11064420817118263 with a P-value of P = 0.08853397846830766

Weight\_pounds

The Pearson Correlation Coefficient for Weight\_pounds is -0.050312258377008784 with a P-value of P = 0.4397693853479999

CPU\_core

The Pearson Correlation Coefficient for CPU\_core is 0.4593977773355115 with a P-value of P = 7.912950127009034e-14

OS

The Pearson Correlation Coefficient for OS is -0.22172980114827384 with a P-value of P = 0.0005696642559246749

GPU

The Pearson Correlation Coefficient for GPU is 0.2882981988881428 with a P-value of P = 6.166949698364282e-06

Category

The Pearson Correlation Coefficient for Category is 0.28624275581264125 with a P-value of P = 7.225696235806733e-06